

This is a companion notebook for the book Deep Learning with Python, Second Edition. For readability, it only contains runnable code blocks and section titles, and omits everything else in the book: text paragraphs, figures, and pseudocode.

If you want to be able to follow what's going on, I recommend reading the notebook side by side with your copy of the book.

This notebook was generated for TensorFlow 2.6.

Deep learning for timeseries

Different kinds of timeseries tasks

A temperature-forecasting example

In []:

```
!wget https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
```

```
!unzip jena_climate_2009_2016.csv.zip
```

Inspecting the data of the Jena weather dataset

In []:

```
import os
```

```
fname = os.path.join("jena_climate_2009_2016.csv")
```

```
with open(fname) as f:
```

```
    data = f.read()
```

```
lines = data.split("\n")
```

```
header = lines[0].split(",")
```

```
lines = lines[1:]
```

```
print(header)
```

```
print(len(lines))
```

Parsing the data

```
In [ ]:
```

```
import numpy as np
```

```
temperature = np.zeros((len(lines),))
```

```
raw_data = np.zeros((len(lines), len(header) - 1))
```

```
for i, line in enumerate(lines):
```

```
    values = [float(x) for x in line.split(",")[1:]]
```

```
    temperature[i] = values[1]
```

```
    raw_data[i, :] = values[:]
```

Plotting the temperature timeseries

```
In [ ]:
```

```
from matplotlib import pyplot as plt
```

```
plt.plot(range(len(temperature)), temperature)
```

Plotting the first 10 days of the temperature timeseries

```
In [ ]:
```

```
plt.plot(range(1440), temperature[:1440])
```

Computing the number of samples we'll use for each data split

```
In [ ]:
```

```
num_train_samples = int(0.5 * len(raw_data))
```

```
num_val_samples = int(0.25 * len(raw_data))  
num_test_samples = len(raw_data) - num_train_samples - num_val_samples  
print("num_train_samples:", num_train_samples)  
print("num_val_samples:", num_val_samples)  
print("num_test_samples:", num_test_samples)
```

Preparing the data

Normalizing the data

In []:

```
mean = raw_data[:num_train_samples].mean(axis=0)  
raw_data -= mean  
std = raw_data[:num_train_samples].std(axis=0)  
raw_data /= std
```

In []:

```
import numpy as np  
from tensorflow import keras  
int_sequence = np.arange(10)  
dummy_dataset = keras.utils.timeseries_dataset_from_array(  
    data=int_sequence[:-3],  
    targets=int_sequence[3:],  
    sequence_length=3,  
    batch_size=2,  
)
```

for inputs, targets in dummy_dataset:

```
for i in range(inputs.shape[0]):  
    print([int(x) for x in inputs[i]], int(targets[i]))
```

Instantiating datasets for training, validation, and testing

In []:

```
sampling_rate = 6
```

```
sequence_length = 120
```

```
delay = sampling_rate * (sequence_length + 24 - 1)
```

```
batch_size = 256
```

```
train_dataset = keras.utils.timeseries_dataset_from_array(  
    raw_data[: -delay],  
    targets=temperature[delay:],  
    sampling_rate=sampling_rate,  
    sequence_length=sequence_length,  
    shuffle=True,  
    batch_size=batch_size,  
    start_index=0,  
    end_index=num_train_samples)
```

```
val_dataset = keras.utils.timeseries_dataset_from_array(  
    raw_data[: -delay],  
    targets=temperature[delay:],  
    sampling_rate=sampling_rate,  
    sequence_length=sequence_length,
```

```
shuffle=True,  
batch_size=batch_size,  
start_index=num_train_samples,  
end_index=num_train_samples + num_val_samples)
```

```
test_dataset = keras.utils.timeseries_dataset_from_array(  
    raw_data[: -delay],  
    targets=temperature[delay:],  
    sampling_rate=sampling_rate,  
    sequence_length=sequence_length,  
    shuffle=True,  
    batch_size=batch_size,  
    start_index=num_train_samples + num_val_samples)
```

Inspecting the output of one of our datasets

In []:

for samples, targets in train_dataset:

```
    print("samples shape:", samples.shape)
```

```
    print("targets shape:", targets.shape)
```

```
    break
```

A common-sense, non-machine-learning baseline

Computing the common-sense baseline MAE

In []:

```
def evaluate_naive_method(dataset):
```

```

total_abs_err = 0.
samples_seen = 0
for samples, targets in dataset:
    preds = samples[:, -1, 1] * std[1] + mean[1]
    total_abs_err += np.sum(np.abs(preds - targets))
    samples_seen += samples.shape[0]
return total_abs_err / samples_seen

```

```
print(f"Validation MAE: {evaluate_naive_method(val_dataset):.2f}")
```

```
print(f"Test MAE: {evaluate_naive_method(test_dataset):.2f}")
```

Let's try a basic machine-learning model

Training and evaluating a densely connected model

In []:

```
from tensorflow import keras
```

```
from tensorflow.keras import layers
```

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
```

```
x = layers.Flatten()(inputs)
```

```
x = layers.Dense(16, activation="relu")(x)
```

```
outputs = layers.Dense(1)(x)
```

```
model = keras.Model(inputs, outputs)
```

```
callbacks = [
```

```
    keras.callbacks.ModelCheckpoint("jena_dense.keras",
```

```

        save_best_only=True)

]

model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])

history = model.fit(train_dataset,

                    epochs=10,

                    validation_data=val_dataset,

                    callbacks=callbacks)

```

```

model = keras.models.load_model("jena_dense.keras")
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")

```

Plotting results

```

In [ ]:

import matplotlib.pyplot as plt

loss = history.history["mae"]

val_loss = history.history["val_mae"]

epochs = range(1, len(loss) + 1)

plt.figure()

plt.plot(epochs, loss, "bo", label="Training MAE")

plt.plot(epochs, val_loss, "b", label="Validation MAE")

plt.title("Training and validation MAE")

plt.legend()

plt.show()

```

Let's try a 1D convolutional model

```

In [ ]:

```

```

inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Conv1D(8, 24, activation="relu")(inputs)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 12, activation="relu")(x)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 6, activation="relu")(x)
x = layers.GlobalAveragePooling1D()(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_conv.keras",
                                    save_best_only=True)
]

model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=10,
                    validation_data=val_dataset,
                    callbacks=callbacks)

```

```

model = keras.models.load_model("jena_conv.keras")
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")

```

A first recurrent baseline

A simple LSTM-based model

In []:

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
```

```
x = layers.LSTM(16)(inputs)
```

```
outputs = layers.Dense(1)(x)
```

```
model = keras.Model(inputs, outputs)
```

```
callbacks = [
```

```
    keras.callbacks.ModelCheckpoint("jena_lstm.keras",  
                                    save_best_only=True)
```

```
]
```

```
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
```

```
history = model.fit(train_dataset,
```

```
                    epochs=10,
```

```
                    validation_data=val_dataset,
```

```
                    callbacks=callbacks)
```

```
model = keras.models.load_model("jena_lstm.keras")
```

```
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")
```

Understanding recurrent neural networks

NumPy implementation of a simple RNN

In []:

```
import numpy as np
```

```
timesteps = 100
```

```
input_features = 32
```

```

output_features = 64

inputs = np.random.random((timesteps, input_features))

state_t = np.zeros((output_features,))

W = np.random.random((output_features, input_features))

U = np.random.random((output_features, output_features))

b = np.random.random((output_features,))

successive_outputs = []

for input_t in inputs:

    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)

    successive_outputs.append(output_t)

    state_t = output_t

final_output_sequence = np.stack(successive_outputs, axis=0)

```

A recurrent layer in Keras

An RNN layer that can process sequences of any length

In []:

```

num_features = 14

inputs = keras.Input(shape=(None, num_features))

outputs = layers.SimpleRNN(16)(inputs)

```

An RNN layer that returns only its last output step

In []:

```

num_features = 14

steps = 120

inputs = keras.Input(shape=(steps, num_features))

```

```
outputs = layers.SimpleRNN(16, return_sequences=False)(inputs)
print(outputs.shape)
```

An RNN layer that returns its full output sequence

In []:

```
num_features = 14
steps = 120
inputs = keras.Input(shape=(steps, num_features))
outputs = layers.SimpleRNN(16, return_sequences=True)(inputs)
print(outputs.shape)
```

Stacking RNN layers

In []:

```
inputs = keras.Input(shape=(steps, num_features))
x = layers.SimpleRNN(16, return_sequences=True)(inputs)
x = layers.SimpleRNN(16, return_sequences=True)(x)
outputs = layers.SimpleRNN(16)(x)
```

Advanced use of recurrent neural networks

Using recurrent dropout to fight overfitting

Training and evaluating a dropout-regularized LSTM

In []:

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(32, recurrent_dropout=0.25)(inputs)
x = layers.Dropout(0.5)(x)
```

```

outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_lstm_dropout.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=50,
                    validation_data=val_dataset,
                    callbacks=callbacks)

```

In []:

```

inputs = keras.Input(shape=(sequence_length, num_features))
x = layers.LSTM(32, recurrent_dropout=0.2, unroll=True)(inputs)

Stacking recurrent layers

Training and evaluating a dropout-regularized, stacked GRU model

```

In []:

```

inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.GRU(32, recurrent_dropout=0.5, return_sequences=True)(inputs)
x = layers.GRU(32, recurrent_dropout=0.5)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_stacked_gru_dropout.keras",
                                    save_best_only=True)
]

model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=50,
                    validation_data=val_dataset,
                    callbacks=callbacks)

model = keras.models.load_model("jena_stacked_gru_dropout.keras")
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")

```

Using bidirectional RNNs

Training and evaluating a bidirectional LSTM

In []:

```

inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Bidirectional(layers.LSTM(16))(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=10,
                    validation_data=val_dataset)

```